

Dask

Parallel and Distributed Computing in Python

April 17, 2018

Seminar Distributed Data Mining
Leiden University

Table of Contents

1. Introduction
2. Collections
3. Task Graphs
4. Workers and Schedulers
5. Task Scheduling
6. Conclusions
7. Dask Demo

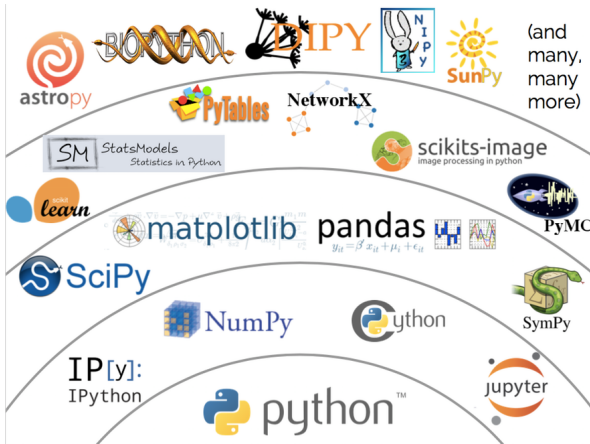
Introduction

What is Dask?

- Relatively new, open source, pure python library (around 2015).
- A project of Continuum Analytics.
- A parallel and distributed computing framework.
- Leverages scientific python ecosystem (numpy, pandas, scikit-learn etc).
- Utilizes block algorithms and task scheduling.

Why Dask?


- Parallelize scientific python ecosystem.



Why Dask? (Cont.)

- Familiar API.
- Scales up to a cluster (Distributed).
- Scales down to a single computer (Parallel).
- Work on data larger than memory.


Task Overview



- Collections create task graphs.
- Scheduler executes graphs on parallel hardware (single machine or cluster).

Collections


Numpy



```
# NumPy code
import numpy as np
x = np.random.random((1000, 1000))
mean = np.mean(x, axis=1)

# Dask.array code
import dask.array as da
x = da.random.random((100000, 100000), chunks = (1000, 1000))
mean = da.mean(x, axis=1)
mean.compute()
```

Pandas



```
# pandas code
import pandas as pd
data = pd.read_csv('sample_data.csv')
data.groupby(data['reviewerID'])['overall'].mean()

# Dask.dataframe code
import dask.dataframe as dd
data = dd.read_csv('sample_data.csv')
mean_per_reviewerID = data.groupby(data['reviewerID'])['overall'].mean()
mean_per_reviewerID.compute()
```

Dask.delayed


- Not every problem is a numpy array, or pandas dataframe.
- What about generic code?
- Dask.delayed mimics for loops, and wraps custom code.

```
def inc(x):  
    return x + 1  
  
def double(x):  
    return x + 2  
  
def add(x, y):  
    return x + y  
  
data = [1, 2, 3, 4, 5]  
  
output = []  
for x in data:  
    a = inc(x)  
    b = double(x)  
    c = add(a, b)  
    output.append(c)  
  
total = sum(output)
```

```
import dask.delayed  
# Dask.delayed  
output = []  
for x in data:  
    a = dask.delayed(inc)(x)  
    b = dask.delayed(double)(x)  
    c = dask.delayed(add)(a, b)  
    output.append(c)  
  
total = dask.delayed(sum)(output)  
  
# No computation has happened yet.  
# The object "total" contains a task graph of the computation.  
  
total.compute()
```

Task Graphs

1D-Array



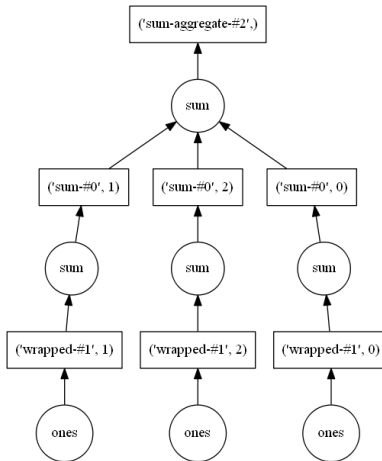
```
In [76]: np.ones((15,))
```

```
Out[76]:
```

```
array([ 1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  
        1.,  1.])
```

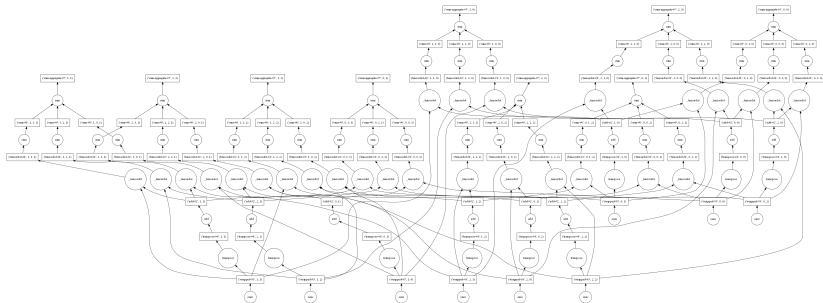
```
In [77]: x = da.ones((15,), chunks = (5,))
```

1D-Array Sum



```
In [82]: x = da.ones((15,), chunks = (5,))  
In [83]: x.sum()
```

ND-Array Matrix Multiplication



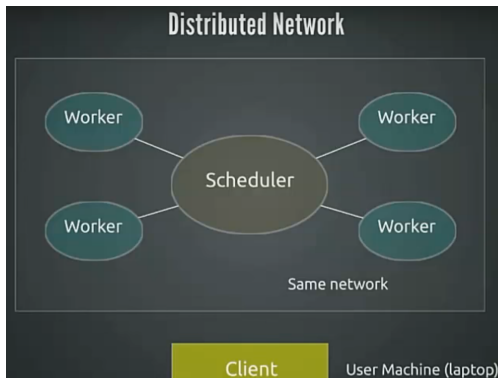
```
In [85]: x = da.ones((15, 15), chunks=(5, 5))
```

```
In [86]: x.dot(x.T + 1)
```

Workers and Schedulers

Distributed Network

- The client uses Dask collections to submit graphs to the scheduler.
- The scheduler distributes jobs to the workers.
- Workers communicate peer-to-peer.



Types of Schedulers

- Dask has two types of schedulers:
 1. Single machine schedulers. Scales only to a single machine.
 2. Distributed scheduler. Scales to a single machine or cluster.
- Single machine schedulers.
 - Synchronous (single-threaded) scheduler. Provides no parallelism, useful for debugging.
 - Threaded scheduler. Provides parallelism for non-Python computations, such as numpy, pandas or any other C/C++ based code.
 - Multiprocessing scheduler. Provides parallelism for generic Python code (strings, lists, dicts etc). However, inter-task data transfer is expensive.
- Distributed Scheduler.
 - Although it's called "Distributed", can be set up either in a local machine or a cluster.
 - Recommended scheduler by Dask developers.

- Workers provide two functionalities
 - Compute tasks as directed by the scheduler.
 - Store results locally and pass them to other workers or clients.
- The scheduler
 - Dynamically assigns tasks to new workers, as they become available.
 - Informs a worker about which peer-workers have the necessary results to compute the assigned task.
 - Deletes previously computed results from memory.

Communication Between Workers and Scheduler

- Two workers named Bob and Alice.
- Scheduler -> Alice: Compute "x <- add(1, 2)"!
 Alice -> Scheduler: I've computed x and am holding on to it!
 Scheduler -> Bob: Compute "y <- add(x, 10)"!
 You will need x. Alice has x.

 Bob -> Alice: Please send me x.
 Alice -> Bob: Sure. x is 3!
 Bob -> Scheduler: I've computed y and am holding on to it!
- Task graph execution (gif):

https://www.dropbox.com/s/ksjoohokgh8hczl/grid_search_schedule.gif?dl=0

Task Scheduling

Journey of a Task (1)

- We follow a single task through user interface, scheduler, worker nodes, and back.
- Task: Addition of two variables `x`, `y` already on the cluster. Pulls the result back to the local process.

```
client = Client()
z = client.submit(add, x, y) # we follow z

print(z.result())
```

Journey of a Task (2)

- Step 1: Client

- `client.submit()` sends the following message to the scheduler

```
{'op': 'update-graph',  
 'tasks': {'z': (add, x, y)},  
 'keys': ['z']}
```

- Step 2: Arrival in the scheduler

- The scheduler creates a graph that shows how to compute z.

- Step 3: Select a worker

1. We consider workers that have either x or y in local memory.
2. Select the worker that would have to gather the least number of bytes to get both x and y locally.
3. Break ties by selecting the least busy worker to gather x or y.

Journey of a Task (3)

- Step 4: Transmit to the worker
 - Task z is placed into a `worker_queue`, in case the worker has not finished a previous task.
 - Information important for the worker is packed in the following message

```
{'op': 'compute',  
 'function': execute_task,  
 'args': ((add, 'x', 'y'),),  
 'who_has': {'x': {(worker_host, port)},  
             'y': {(worker_host, port), (worker_host, port)}},  
 'key': 'z'}
```

- Step 5: Execute on the worker
 - Worker unpacks the message.
 - Collects x or y from other workers (`who_has` key in dictionary).
 - Computes `add(x, y)`.
 - Holds on to the result.
 - Communicates to the scheduler the number of bytes of the result.

Journey of Task (4)

- Step 6: Scheduler aftermath
 - Scheduler sends a new task, if available, to the worker.
 - Scheduler informs workers to delete x , y if no longer needed.
- Step 7: Gather
 - Scheduler informs worker to return task z to the user.
- Step 8: Garbage Collection
 - If no computations depend on z , the scheduler removes elements of task z from its state.

Conclusions

Conclusions (and Some Weak Points)

- Dask scales computations to a single machine or cluster.
- Handles data larger than memory.
- Collections create task graphs.
- Schedulers execute task graphs on parallel hardware.
- Familiar API, easy to get started.
- Some weak points.
 - Does not always work smoothly (memory leaks).
 - Some computations might be slower with Dask, i.e pandas groupby, merge.

Additional Resources

- Website and documentation:
<https://dask.pydata.org/en/latest/>
- Documentation for distributed scheduler:
<https://distributed.readthedocs.io/en/latest/>
- Source code: <https://github.com/dask>
- Examples: <https://github.com/dask/dask-examples>
- Matthew Rocklin's blog: <http://matthewrocklin.com/blog/>
- Online course on DataCamp: <https://www.datacamp.com/courses/parallel-computing-with-dask>
- Conference talks on YouTube.

Dask Demo

Thank you for your attention!

Questions?